New York University Tandon School of Engineering
Computer Science and Engineering

CS-GY 6763: Midterm Practice.

Logistics

- Exam will be held in class on **Friday, 3/14** starting at 2:00pm sharp. Please arrive on time!
- You will have **1 hour, 15 minutes** to answer a variety of short answer and longer form questions.
- You can bring a one page sheet of paper (two-sided if you want) with notes, theorems, etc. written down for reference.
- Scrap paper will be provided and can be used to write solutions if extra space is needed.
- I will be in the room to answer any questions.

Concepts to Know

Random variables and concentration.

- Linearity of expectation and variance.
- Indicator random variables and how to use them.
- Markov's inequality, Chebyshev's inequality (ideally should know from memory so you can apply quickly).
- Union bound (should know from memory).
- Chernoff and Bernstein bounds (don't need to memorize the exact bounds, but can apply if given).
- General idea of law of large numbers and central limit theorem.
- The probability that a normal random variables $\mathcal{N}(0, \sigma^2)$ falls further than $k\sigma$ away from its expectation is $\leq O(e^{-k^2/2})$.

Hashing, Dimensionality Reduction, High Dimensional Vectors

- Random hash functions.
- Random hashing for frequency estimation.
- Random hashing for distinct elements estimation.
- MinHash for Jaccard similarity estimation.
- Locality sensitive hash functions.
- MinHash and SimHash for Jaccard Similarity and Cosine Similarity.
- Adjusting false positive rate and false negative rate in an LSH scheme.
- Statement of Johnson-Lindenstrauss lemma (know from memory).
- Statement of *distributional* JL lemma and how it can be used to prove JL.

High dimensional geometry

- How to draw a random unit vector from the sphere in d dimensions (draw $\mathbf{x}$ with all entries i.i.d. $\mathcal{N}(0,1)$ and normalize it).
- How does $\|x-y\|_2^2$ relate to $\langle x, y\rangle$ if x and y are unit vectors?
- How many mutually orthogonal unit vectors are there in d dimensions?
- There are $2^{\theta(\epsilon^2 d)}$ nearly orthogonal unit vectors in d dimensions (with $\langle x, y\rangle \leq \epsilon$). Know roughly how prove this fact using the *probabilistic method*, which required a an exponential *concentration inequality* + *union bound*.
- Know how to prove that all but an $2^{\theta(-\epsilon d)}$ fraction of a balls volume in d dimensions lies in a spherical shell of width ϵ near its surface.
- The surface area/volume ratio *increases* in high dimensions.
- The cube volume/ball volume ratio *increases* in high dimensions.

Practice Problems

1. Show that for any random variable X, $\mathbb{E}[X^2] \geq \mathbb{E}[X]^2$.
2. Show that for independent X and Y with $\mathbb{E}[X] = \mathbb{E}[Y] = 0$, $\text{var}[X \cdot Y] = \text{var}[X] \cdot \text{var}[Y]$.
3. Given a random variable X, can we conclude that $\mathbb{E}[1/X] = 1/E[X]$? If so, prove this. If not, give an example where the equality does not hold.
4. Indicate whether each of the following statements is **always** true, **sometimes** true, or **never** true. Provide a short justification for your choice.
   (a) $\Pr[X = s \text{ and } Y = t] > \Pr[X = s]$. ALWAYS SOMETIMES NEVER
   (b) $\Pr[X = s \text{ or } Y = t] \leq \Pr[X = s] + \Pr[Y = t]$. ALWAYS SOMETIMES NEVER
   (c) $\Pr[X = s \text{ and } Y = t] = \Pr[X = s] \cdot \Pr[Y = t]$. ALWAYS SOMETIMES NEVER
5. Let $\mathbf{\Pi}$ be a random Johnson-Lindenstrauss matrix (e.g. scaled random Gaussians) with $O(\log(1/\delta)/\epsilon^2)$ rows. With probability $(1-\delta)$,
$$\min_{\mathbf{x}} \|\mathbf{\Pi A x} - \mathbf{\Pi b}\|_2^2 \leq (1+\epsilon) \min_{\mathbf{x}} \|\mathbf{A x} - \mathbf{b}\|_2^2.$$
   Is the statement true: ALWAYS SOMETIMES NEVER
6. Assume there are 1000 registered users on your site $u_1, \ldots, u_{1000}$, and in a given day, each user visits the site with some probability p_i. The event that any user visits the site is independent of what the other users do. Assume that $\sum_{i=1}^{1000} p_i = 500$.
   (a) Let X be the number of users that visit the site on the given day. What is $E[X]$?
   (b) Apply a Chernoff bound to show that $Pr[X \geq 600] \leq .01$.
   (c) Apply Markov's inequality and Chebyshev's inequality to bound the same probability. How do they compare?
7. Give an example of a random variable and a deviation t where Markov's inequality gives a tighter upper bound than Chebyshev's inequality.
8. Suppose there is some unknown vector $\boldsymbol{\mu} \in \mathbb{R}^d$. We receive noise perturbed random samples of the form $\mathbf{Y}_1 = \boldsymbol{\mu} + \mathbf{X}_1, \ldots, \mathbf{Y}_k = \boldsymbol{\mu} + \mathbf{X}_k$ where each $\mathbf{X}_i$ is a random vector with each of its entries distributed as an independent random normal $\mathcal{N}(0,1)$. From our samples $\mathbf{Y}_1, \ldots, \mathbf{Y}_k$ we hope to estimate $\boldsymbol{\mu}$ by $\tilde{\boldsymbol{\mu}} = \frac{1}{k}\sum_{i=1}^{k} \mathbf{Y}_i$.

(a) Prove that with $k = O(\log d/\epsilon^2)$ samples, $\max_{i=1,\ldots,d} |\boldsymbol{\mu}_i - \tilde{\boldsymbol{\mu}}_{\boldsymbol{i}}| \leq \epsilon$ with probability 9/10.

(b) Prove that with $k = O(d \log d/\epsilon)$ samples, $\|\boldsymbol{\mu} - \tilde{\boldsymbol{\mu}}\|_2 \leq \epsilon$ with probability 9/10.

9. For two length d binary vectors $\mathbf{q}, \mathbf{y} \in \{0,1\}^d$, consider the hamming similarity:

$$s(\mathbf{q}, \mathbf{y}) = 1 - \frac{\|\mathbf{q} - \mathbf{y}\|_0}{d}.$$

Recall that $\|\mathbf{q}-\mathbf{y}\|_0 = \sum_{i=1}^{d} \mathbb{1}[\mathbf{q}_i \neq \mathbf{y}_i]$. Construct a function h as follows: define the random function $c : \{0,1\}^d \to \{0,1\}$ as $c(\mathbf{x}) = \mathbf{x}[j]$, where j is a uniform random integer in $\{1,\ldots,d\}$. Then, let g be a uniform random hash function from $\{0,1\} \to \{1,\ldots,m\}$. Finally, let:

$$h(\mathbf{x}) = g(c(\mathbf{x})).$$

Prove that h is a locality sensitive hash function for hamming similarity.